CS4243
Computer Vision & Pattern Recognition

AY 2023/24

Lab Session 1

NUS National University of Singapore | School of Computing

Teaching Assistants

Lingdong Kong

lingdong.kong@u.nus.edu

Xiao Cao

xiaocao@u.nus.edu

Arrangement

- Part 1 - Quick Recap from the Lecture (~20 min)
- Part 2 - Lab Tutorial (~30 min)
- Break (10 min)
- Part 3 - Lab Solution (~30 min)

Lab Materials

- GitHub Repo: https://github.com/ldkong1205/cs4243_lab
- Slides
- Notebook & Solution
- Other Materials (image, media, etc.)

Canvas Chat

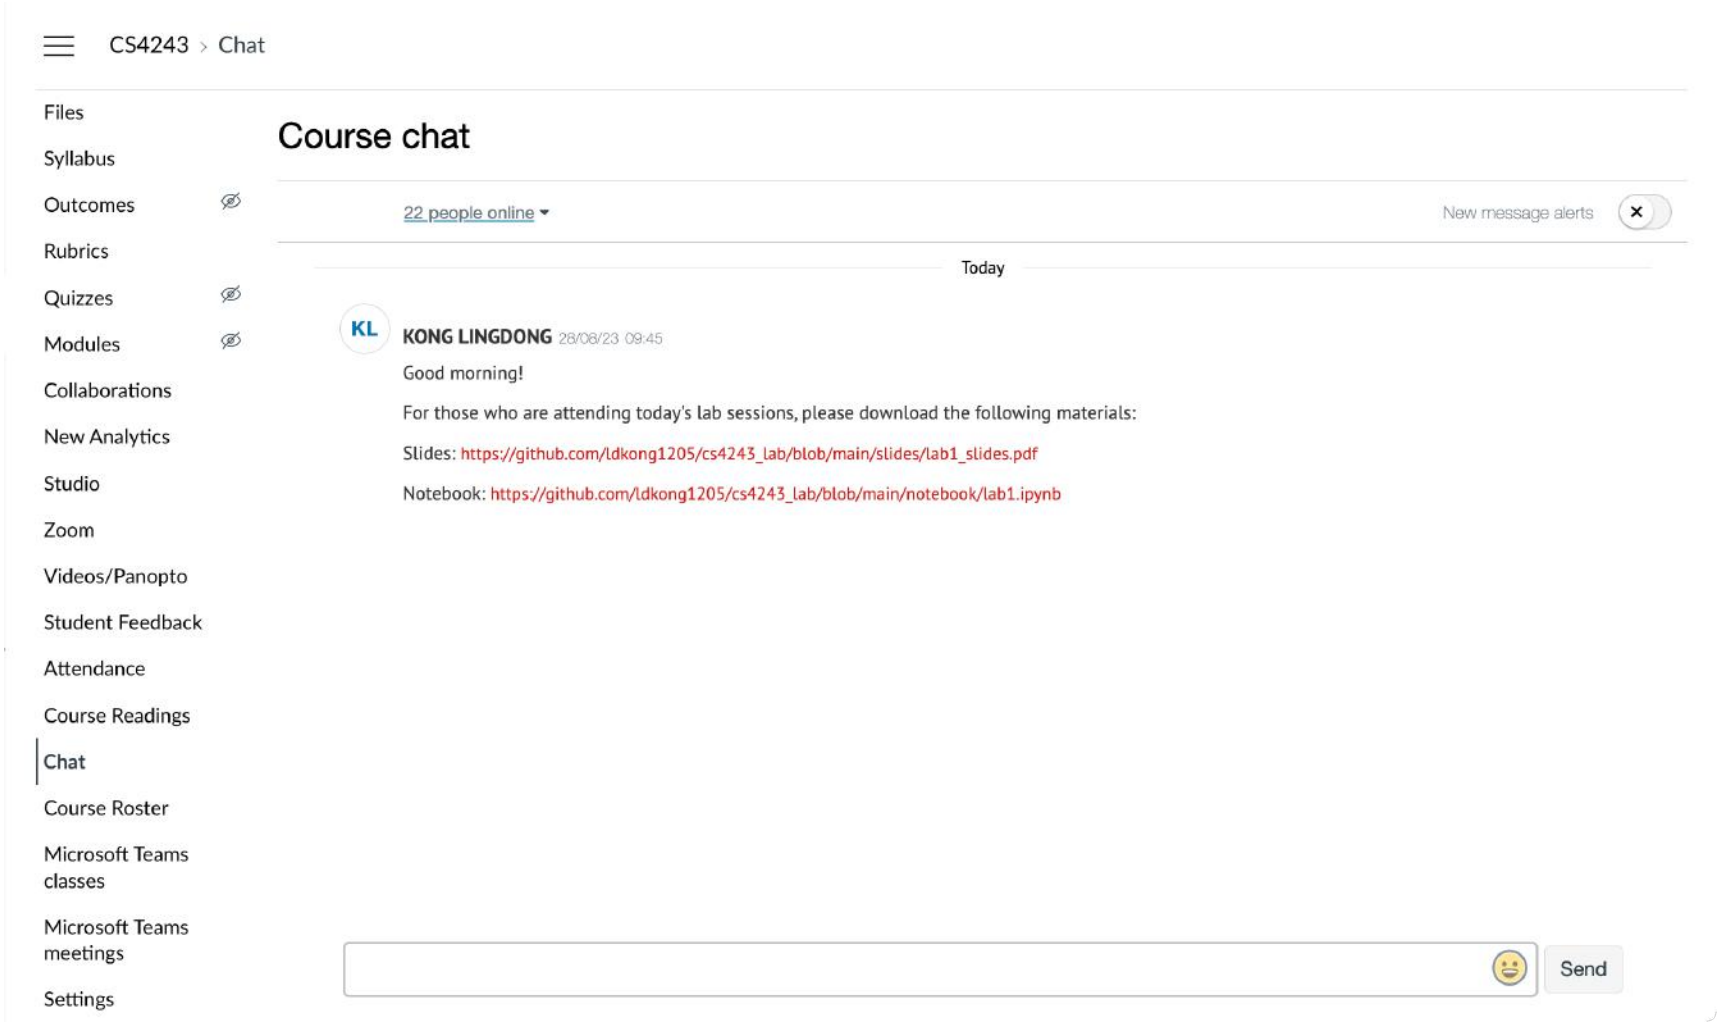
CS4243 › Chat
Files
Syllabus
Outcomes
Rubrics
Quizzes
Modules
Collaborations
New Analytics
Studio
Zoom
Videos/Panopto
Student Feedback
Attendance
Course Readings
Chat
Course Roster
Microsoft Teams classes
Microsoft Teams meetings
Settings
Course chat
22 people online
New message alerts
Today
KL
KONG LINGDONG 28/08/23 09:45
Good morning!
For those who are attending today's lab sessions, please download the following materials:
Slides: https://github.com/ldkong1205/cs4243_lab/blob/main/slides/lab1_slides.pdf
Notebook: https://github.com/ldkong1205/cs4243_lab/blob/main/notebook/lab1.ipynb
Send

Lesson 1

Introduction and Fundamentals

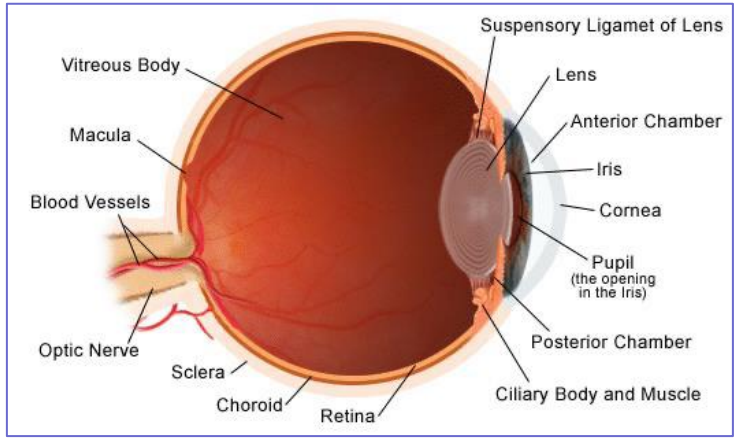

What is this?

A Color Image

- An M x N x 3 matrix
- Each pixel is an 8-bit number between 0 and 255
- RED, GREEN, BLUE
- M x N
  = resolution
  = number of pixels in columns / rows

Caution: OpenCV uses BGR format

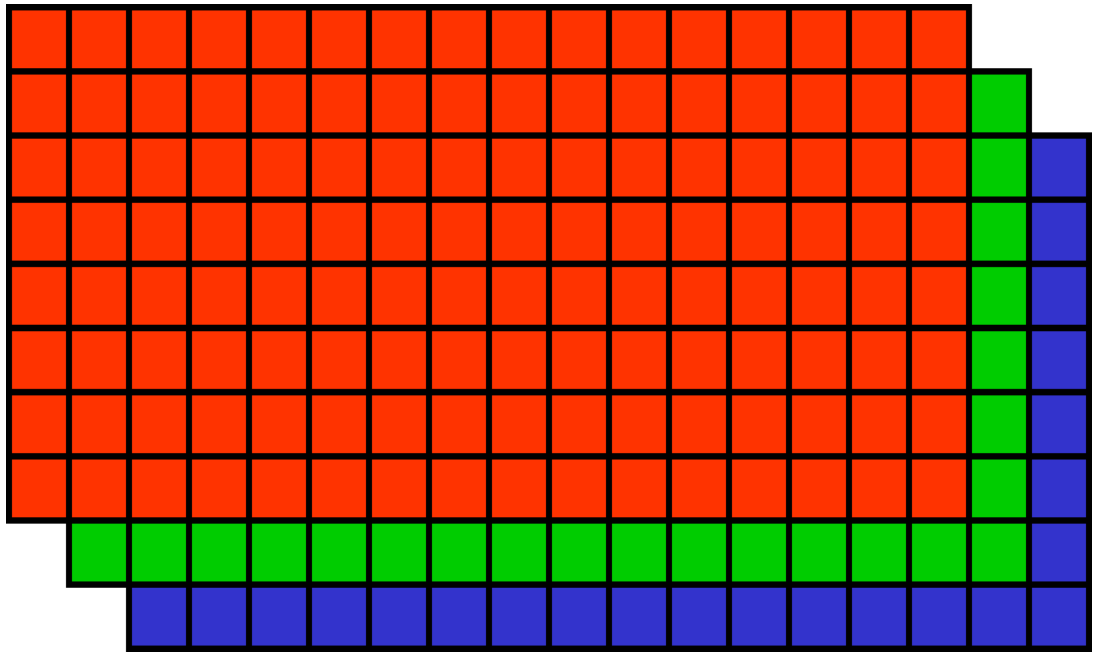

Operations

Anything Can be applied on your Digital Image or Video

Logical Operations	Transforms	Statistical Operations	Geometrical Operations	Mathematical Operations
Morphology, Coding,	Fourier, Walsh, PCA	Histograms, Correlation, Max/Min	Affine Transforms	Filtering, *.**

Histogram

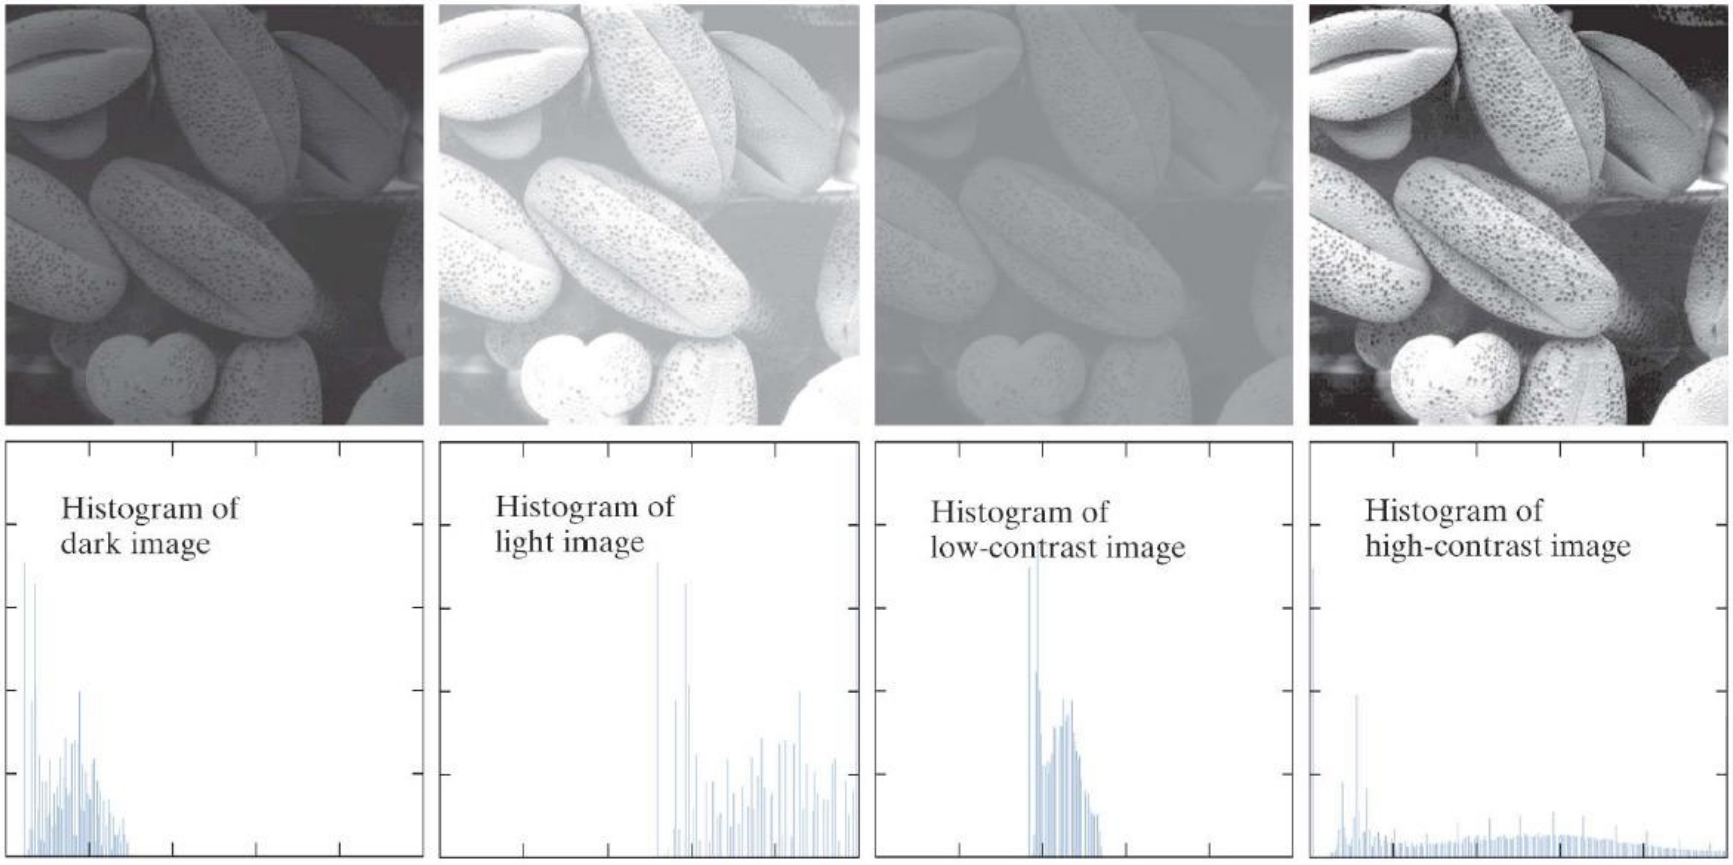

FIGURE 3.16
Four image types and their corresponding histograms. (a) dark; (b) light; (c) low contrast; (d) high contrast. The horizontal axis of the histograms are values of r_k and the vertical axis are values of $p(r_k)$.

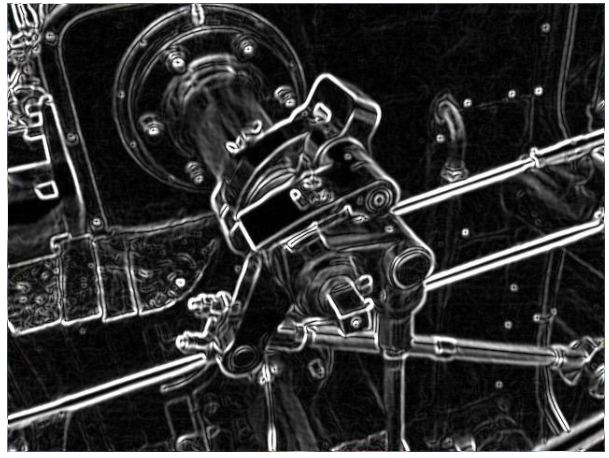

Lesson 2

Local Operations, Filtering, and Convolution

- Where we see how we can filter and image
- Every system in this big universe is a filter

Local Operations

- An image in the spatial domain
- A local operator uses the pixel and its neighbours to compute the new value of that pixel.
- E.g., replace each pixel f(x,y), with the average of that and its 8 neighbours.
- $g(x,y) = T[\ f(x \pm \Delta x, y \pm \Delta y)\ ]$

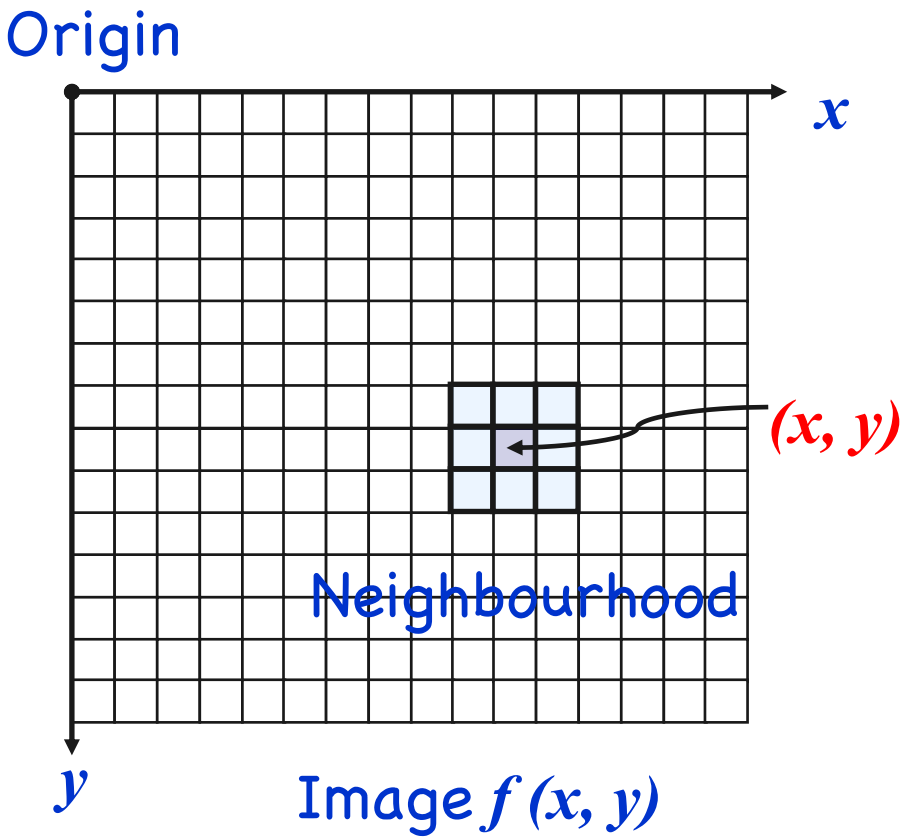

Local Operations

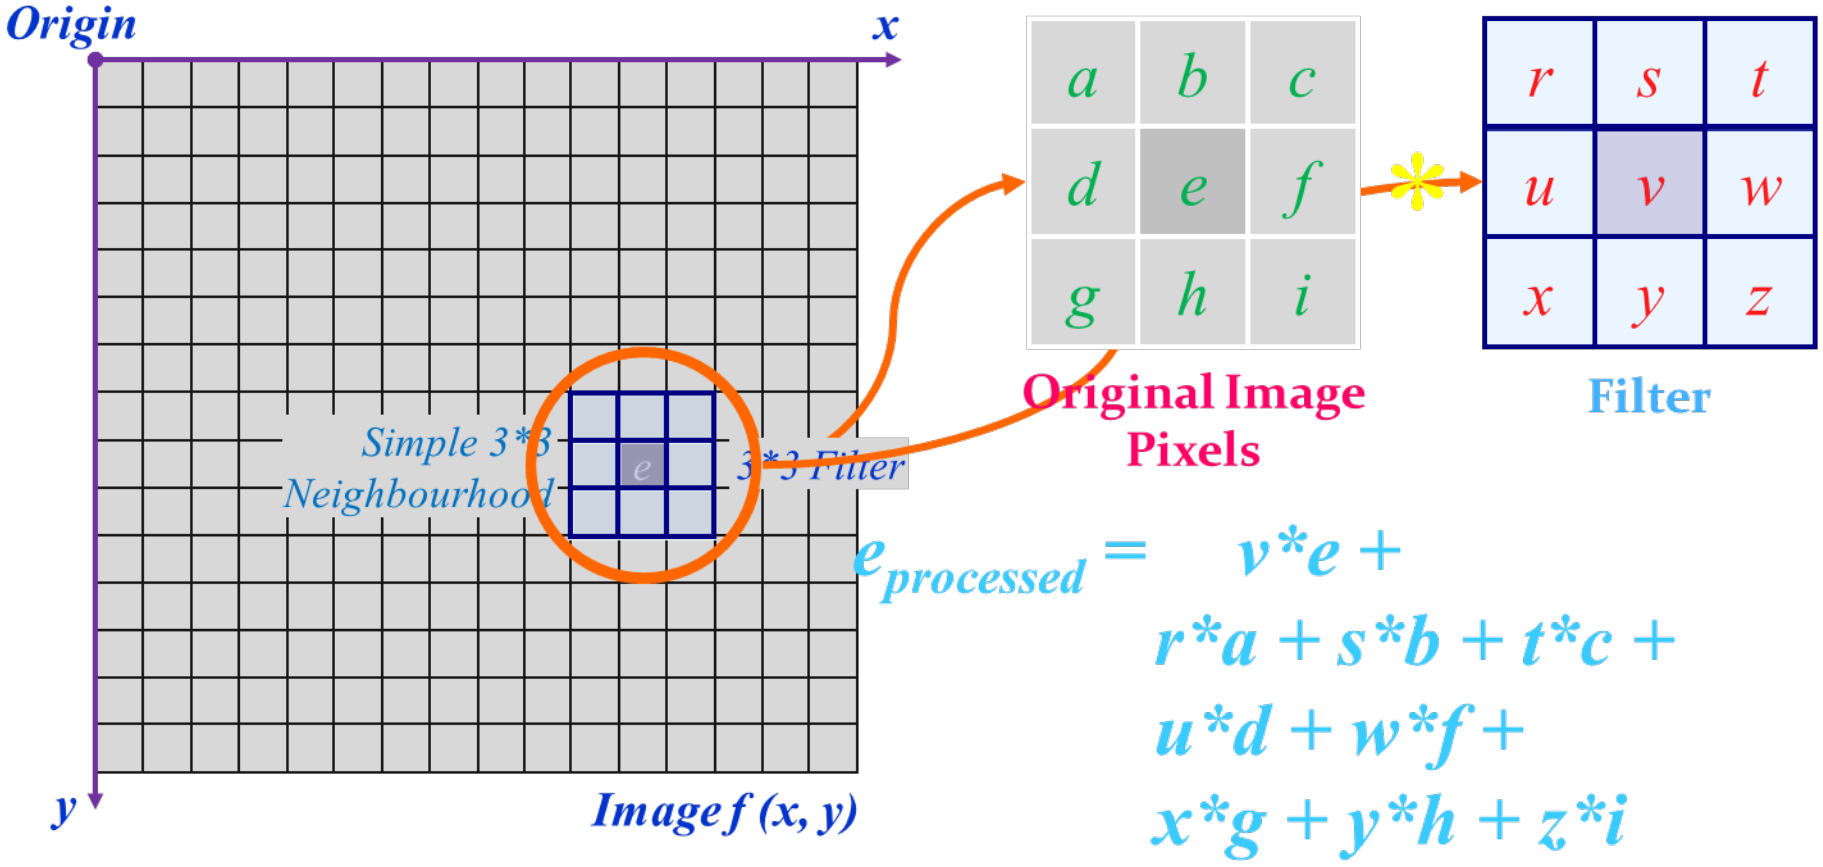

The **filter** would slide over the image, shift or stride is 1 column and then 1 row

Low-Pass Filters

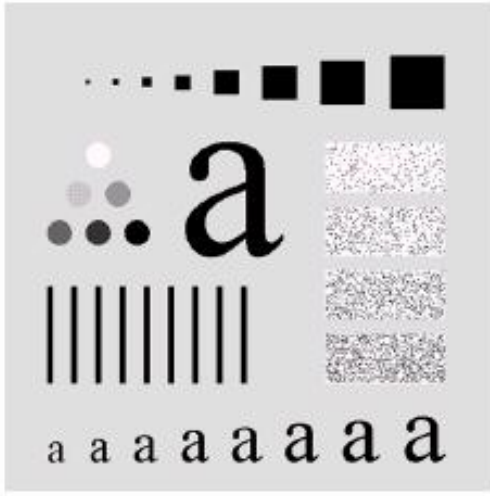

Original image

Filtered with 3x3

Filtered with 5x5

Filtered with 9x9

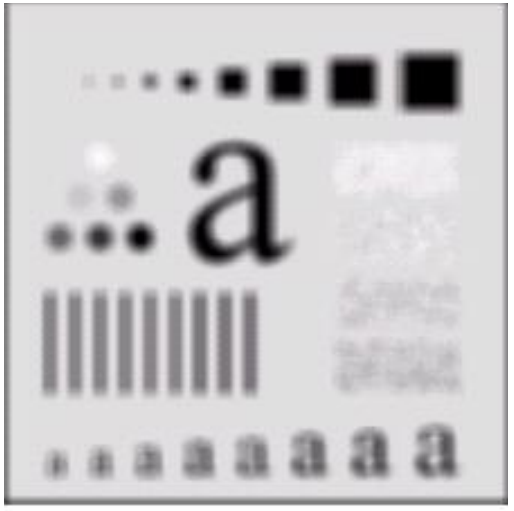

Filtered with 15x15

Filtered with 35x35

High-Pass Filters

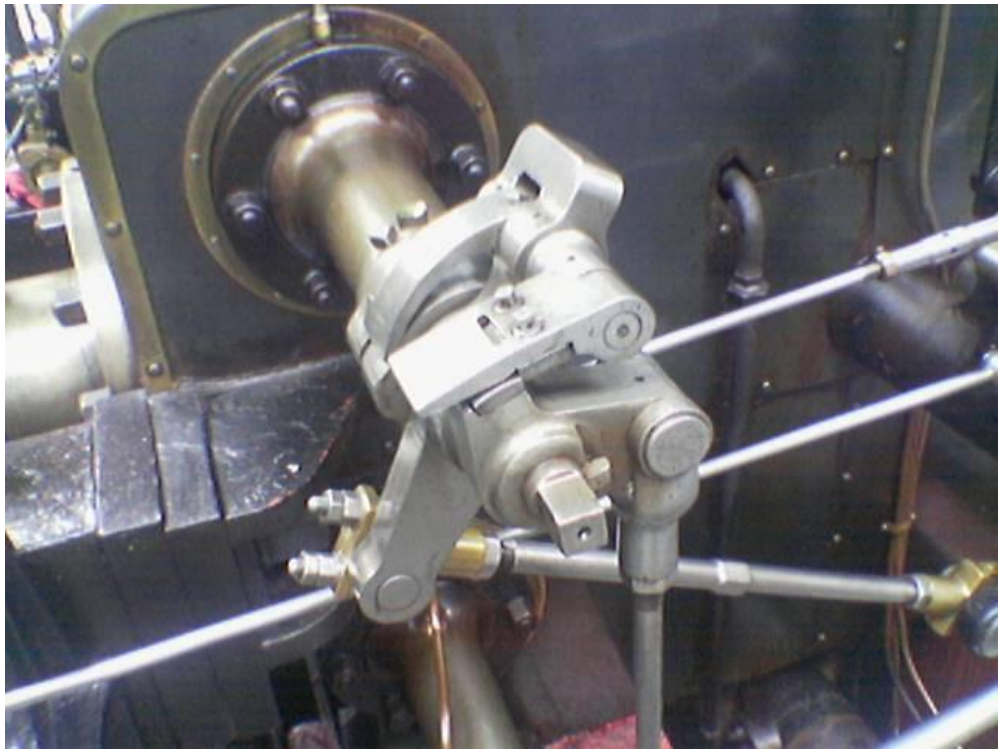

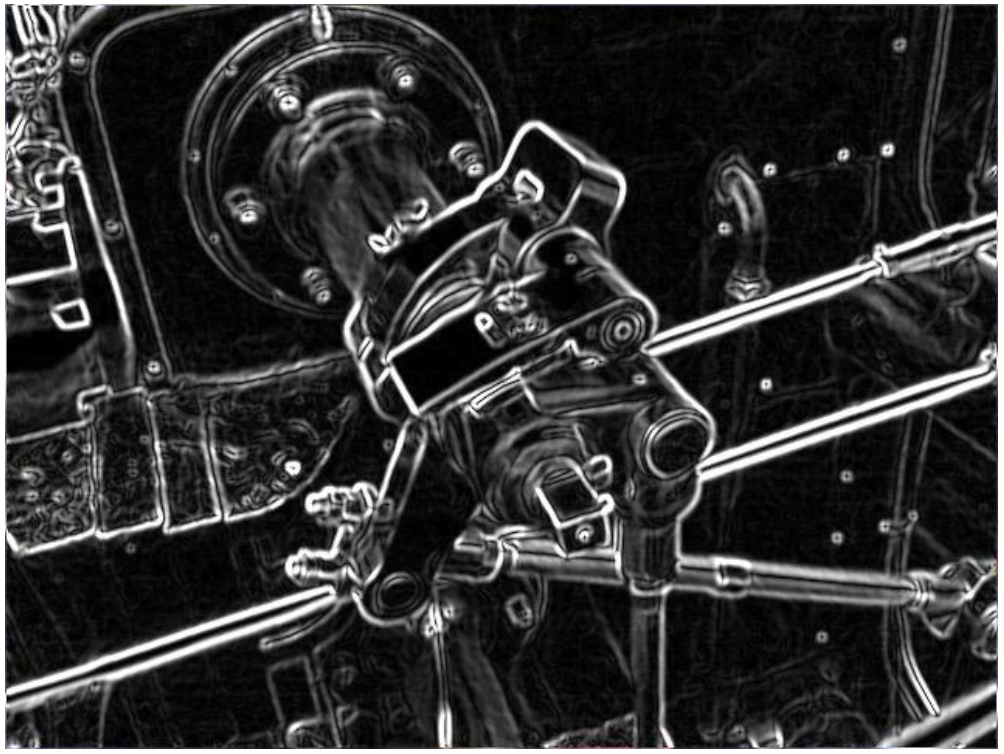

Discrete 1st derivative or **gradient** of an image.

An image and result of edge detection using the Sobel algorithm.

Median Filters

- In a small n x n patch/neighborhood
- Sort the pixels and find the median, *M*
- Replace the central pixel of your patch with *M*
- Slide one pixel to the next neighborhood

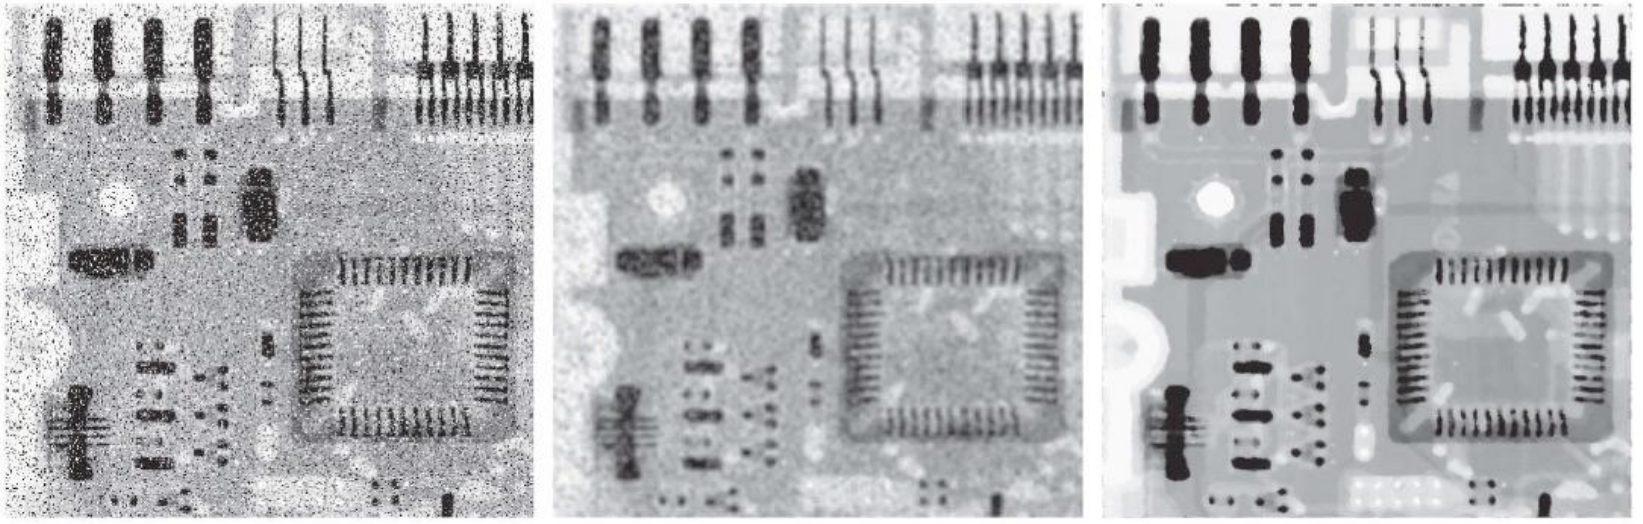

a b c

FIGURE 3.49
(a) X-ray image of a circuit board, corrupted by salt-and-pepper noise. (b) Noise reduction using a 19×19 Gaussian lowpass filter kernel with $\sigma = 3$. (c) Noise reduction using a 7×7 median filter. (Original image courtesy of Mr. Joseph E. Pascente, Lixi, Inc.)

Convolution, 2D

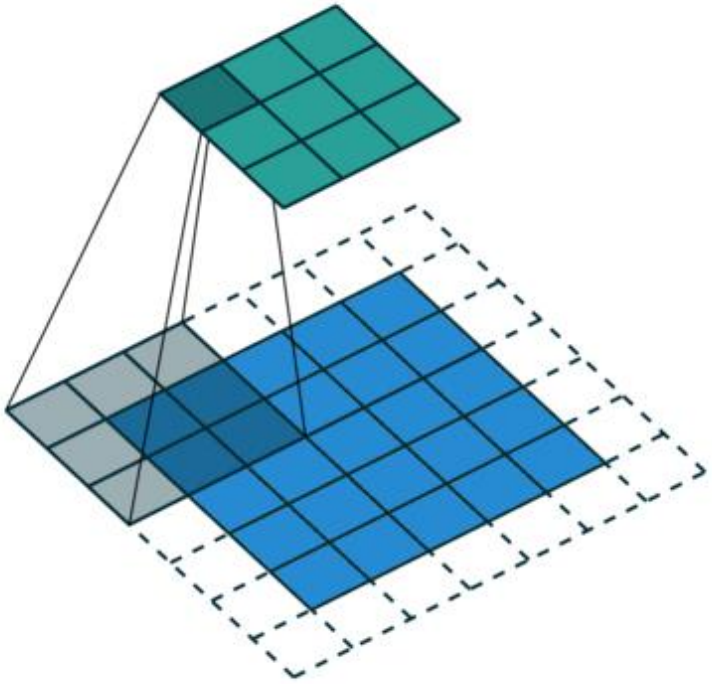

convolution with stride = 1

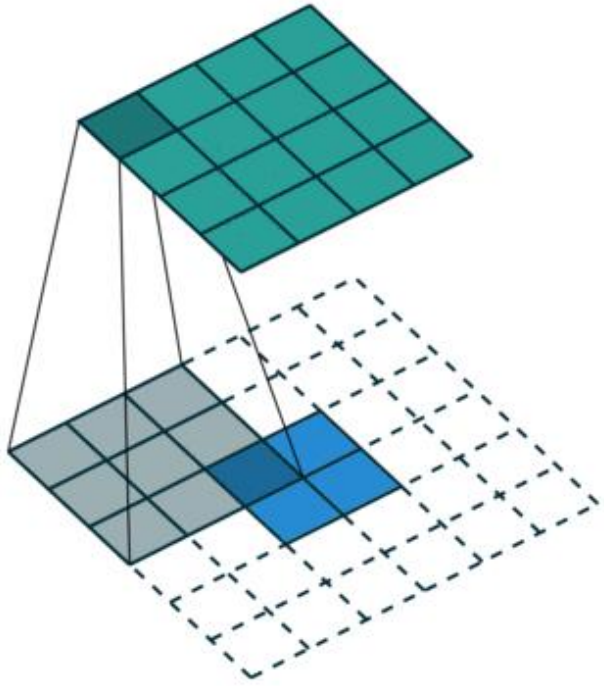

convolution with stride = 2

Convolutional Neural Nets (CNNs)

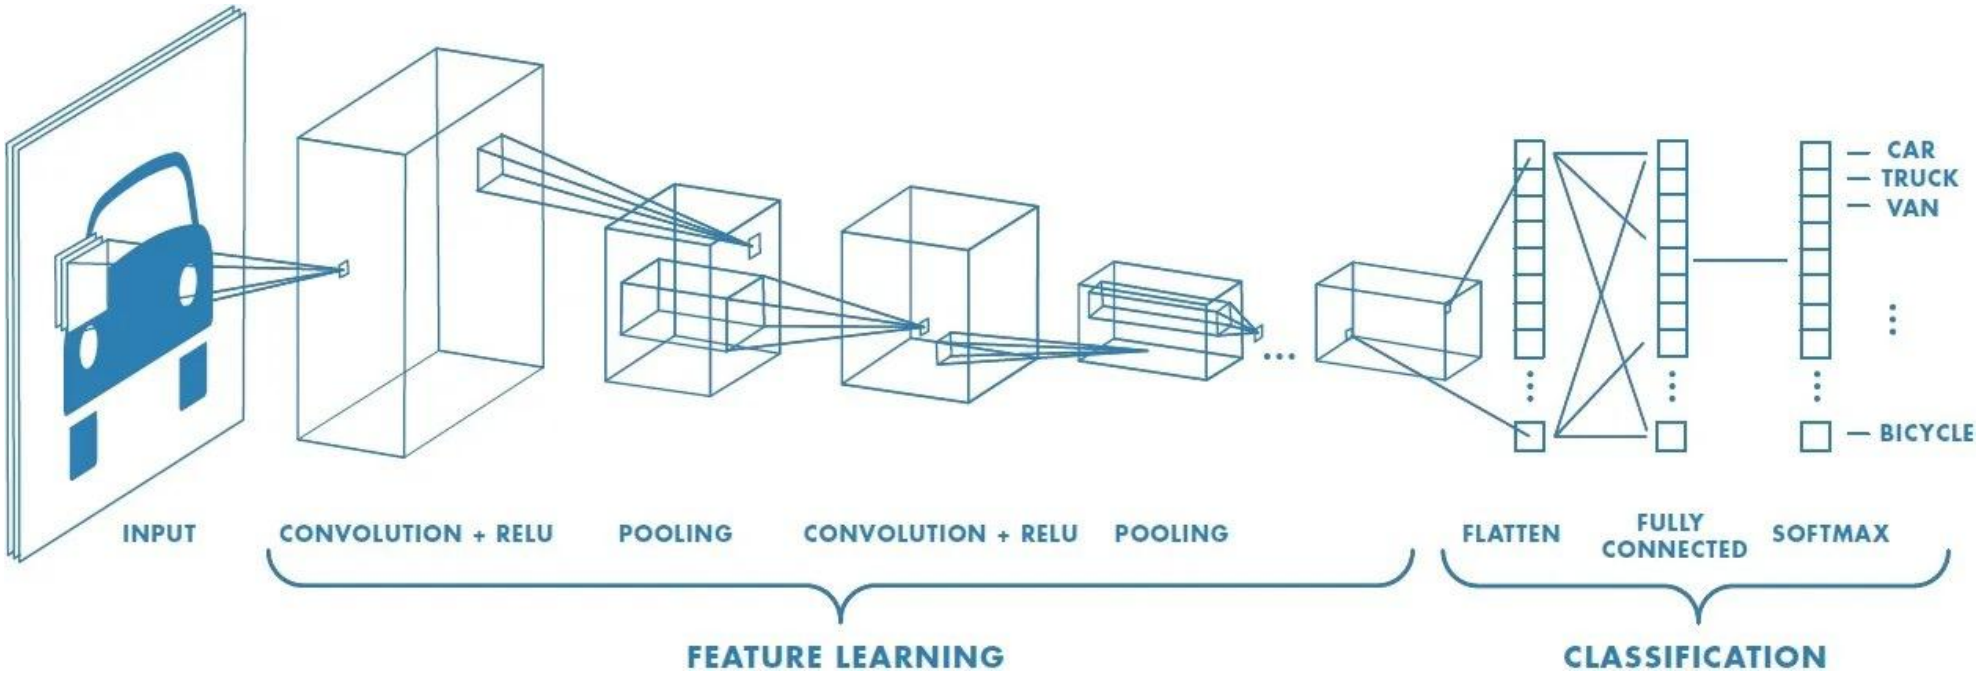

Lab Session 1

Background & Global Operations

NUS National University of Singapore | School of Computing